

WSC

Strong scaling - adding more machines make algorithm faster on a given data

Weak Scaling - adding more machines lets you process more data in the same amount of time

Data Level Parallelism - independent data sets with independent processing.

Big Idea - instead of using expensive super computers use 10,000 - 100,000 cheaper servers & networks

- emphasize cost efficiency
- attention to power: distribution & cooling
- (relatively) homogenous hardware/software, which is good for error handling and scaling

Quick Facts:

- CAPEX - cost to buy equipment (buy servers)
- OPEX - cost to run equipment (electricity used)
- servers cost most \$\$\$ with replacement every 3 years
- building cooling increase PUE
- a lot of servers sit idle (most servers are at 10% - 50% max load) and waste energy
- ultimate goal is energy proportionality - $\frac{\text{\%peak load}}{\text{\%peak energy}}$

Power Usage Effectiveness (PUE) = $\frac{\text{Total Building Power}}{\text{IT Power}}$
can never be less than 1

tells us nothing about the efficiency of our computing hardware, only the relative efficiency of our IT equip vs. non IT equip.

MapReduce

Map - slice data into shards, distribute to workers, compute sub-problem

map(inKey, inVal) → list(outKey, intermedVal)

Combiner - compress the output of a single mapper to save on bandwidth and to distribute work more evenly
apply reduce operation distributed across map tasks, producing set of intermediate pairs

Reducer - collect and combine subproblems, combining all intermediate values for a particular key to produce a set of merged output values

all mappers must finish before any reducer can start

Data Shuffle/Sort (taken care of by framework) merges all intermediate values with same key from mapper output to produce key value pairs of

(intermedKey, list[all associated values])

reduce(outKey, list(intermedVal)) → list(outVal)

C Programming Language

Variable Type	Keyword	Bytes Required	Range
Character	char	1	-128 to 127
Unsigned character	unsigned char	1	0 to 255
Integer	int	2	-32768 to 32767
Short Integer	short int	2	-32768 to 32767
Long Integer	long int	4	-2,147,483,648 to 2,147,438,647
Unsigned Integer	unsigned int	2	0 to 65535
Unsigned Short Integer	unsigned short int	2	0 to 65535
Unsigned Long Integer	unsigned long int	4	0 to 4,294,967,295
Float	float	4	1.2E-38 to
Double	double	8	2.2E-308 to
Long Double	long double	10	3.4E-4932 to 1.1E+4932

Category	Operator	Associativity
Postfix	<code>[] -> . ++ --</code>	Left to right
Unary	<code>+ - ! ~ ++ -- (type) * & sizeof</code>	Right to left
Multiplicative	<code>* / %</code>	Left to right
Additive	<code>+ -</code>	Left to right
Shift	<code><< >></code>	Left to right
Relational	<code>< <= > >=</code>	Left to right
Equality	<code>== !=</code>	Left to right
Bitwise AND	<code>&</code>	Left to right
Bitwise XOR	<code>^</code>	Left to right
Bitwise OR	<code> </code>	Left to right
Logical AND	<code>&&</code>	Left to right
Logical OR	<code> </code>	Left to right
Conditional	<code>?:</code>	Right to left
Assignment	<code>= += -= *= /= %= >>= <<= &= ^= =</code>	Right to left
Comma	<code>,</code>	Left to right

C arrays are almost identical to pointers

array variable is a pointer to the first (0^{th}) element

arr[0] = *arr

arr[2] = *(arr+2)

must pass array & size because array in C does not know its own length

arr[i] = *(arr + i)

MIPS

registers are 32 bits wides (this quantity is called a *word*), 4 bytes

memory addresses are really in bytes, not words word

addresses are 4 bytes apart

jal - jumps to address and simultaneously saves address of

following instruction in register \$ra

bible of MIPS - the green MIPS sheet!

data transfer object: constant(register used to access memory)

→ sum the memory address

address in a word match the address of 1 of the 4 bytes within that word

addresses of sequential words differ by 4 bytes

MIPS words must start at addresses that are multiples of 4

e.g. *beq \$t0, \$0, 2* ← means 2 instructions, each instruction is 4 bytes (32 bits)

C Memory Model

- **Stack**
 - local variables, broken into stack frames, grows down
- **Heap**
 - stores data for which space was allocated with malloc(), grows upwards
- **Static**
 - global, static variables, & string literals
 - fixed size during execution
- **Code**
 - contains instructions, fixed size during exec

High Level → Compiler → Assembly Language → Assembler → Binary Instructions

Below the Program

Operating System

interfaces between user's application & hardware

- handles input/output operations
- allocates memory & storage
- provides sharing of computer among multiple applications

Compiler

a program that translates high level language statements into assembly language statements, symbolic representations of machine instructions

Assembler

input - assembly language code

output - object code

reads & uses directives, replaces pseudo instructions, produces machine language, & creates object files

assembler directions:

- .text - subsequent items put in user text segment
- .data - subsequent items put in user data segment
- .globl sum - declares sum global and can be referenced from other files
- .ascii str - store the string str in memory and null terminate it
- .word w₁ - store a s32 bit quantity in successive memory words

Object Module

assembler or compiler translates program into machine instructions

provides information for building a complete program from the pieces:

- header: contents of object module
- text: translated instructions
- static data segment: data allocated for life of program
- relocation info: for contents that depend on absolute location of program
- symbol table: global definitions and external references
- debug info: for associating with source code

Linker

produces executable image

- merges segments
- resolves labels
- patches location dependent and external references

Loader

- reads header to determine segment sizes
- creates virtual address space
- copy text and initialized data into memory
- sets up arguments on stack
- initializes registers
- jumpstarts routine

(interpreters are slower than compiled programs, but useful for debugging)

Underlying Hardware

Five Components

- input: mechanism by which comp is fed info
- output: mechanism by which computer conveys result of computation
- memory - storage area in which programs are kept when running and area containing data necessary by program to run
- data path - component of the processor that performs arithmetic operations
- control - component of the processor that commands the data path, memory, & I/O devices according to the program instructions

Floating Point

single precision: 1 sign bit (31), 8 exponent bits (30-23), 23

fraction bits (22-0)

$$FP = (-1)^S \times (1 + F) \times 2^{E - bias(127)}$$

special values:

- \pm Zero: E = 0, F = 0
- NaN: E = 255 (all 1's), F \neq 0
- $\pm\infty$: E = 255 (all 1's), F = 0
- Denormalized: E = 0, F \neq 0

underflow: negative exponent is too large to fit in exponent field

overflow: exponent is too large to be represented in exponent field

Performance

throughput/bandwidth = total amount of work done in a given time

to maximize performance, minimize response/execution time

$$\text{Performance} = \frac{1}{\text{executiontime}} \\ \frac{\text{performance}_x}{\text{performance}_y} = \frac{\text{executiontime}_y}{\text{executiontime}_x} = n, \rightarrow x \text{ is } n \text{ times faster than } y$$

response time - time to complete a task

CPU execution time = time CPU spends computing a task not including time spent waiting for I/O

User CPU time - CPU time spent in the program itself

System CPU time - CPU time spent in the operating system performing tasks on behalf of the program

clock cycles - discrete time intervals determining when events take place in the hardware

clock period - time for a complete clock cycle

clock rate - inverse of the clock period

Performance Equations

$$\text{CPU execution time} = \text{CPU clock cycles for a program} \times \text{Clock Cycle time} = \frac{\text{CPU clock cycles for a program}}{\text{clockrate}(\frac{1}{\text{clockcycle time}})}$$

CPU clock cycles = instructions for a program \times average clock cycle per instruction

CPI (clock cycles per instruction) - average number of clock cycles per instruction for a program, an average of all instructions executed in a program.

$$\text{CPU time} = \text{instruction count} \times \text{CPU} \times \text{clock cycle time} = \frac{\text{instructioncount} \times \text{CPI}}{\text{clockrate}}$$

$$\text{CPI} = \frac{\text{CPU clock cycles}}{\text{instructioncount}}$$

$$\text{Time} = \frac{\text{seconds}}{\text{program}} = \frac{\text{instructions}}{\text{program}} \times \frac{\text{clockcycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{clockcycle}}$$

Cache

AMAT (avg. mem. access time) = hit time + miss rate \times miss time

local miss rate = fraction of references to one level of cache that miss, e.g. local miss rate of L2 = $\frac{L2 \text{ misses}}{L1 \text{ misses}}$

global miss rate = fraction of references that miss in all levels of a multi-level cache

Tag: tells us where it came from in memory.

Tag bits = address size - index - offset

Index: which row/block to look at in cache.

Index bits = $\log_2(\frac{\text{cachesize}}{\text{blocksize}})$, in other words log of the number of blocks

Offset: which column of cache to look at.

Offset bits = $\log_2(\text{block size (bytes)})$

bits per row = valid bit + dirty bit (write back cache only) + data (block size) + tag

cache size = $2^{\text{offsetbits} + \text{indexbits}}$

number of blocks = cache size \div block size

tag bits = $\log_2(\frac{\text{memorysize}}{\text{cachesize}})$

$\text{CPI}_{\text{stall}} = \text{CPI}_{\text{ideal}} + \text{data miss cycles} + \text{instruction miss cycles}$

$\text{CPI}_{\text{stall}} = \text{CPI}_{\text{ideal}} + \text{L1 instruction miss cycles} + \text{L1 data miss cycles}$

Types of Instructions on Types of Data

SISD: Single instruction, single data stream, e.g. Intel Pentium 4

MISD: Multiple instruction, single data stream, e.g. none

MIMD: Multiple instruction, multiple data stream, e.g. Intel Xeon e5345

SIMD: Single instruction, multiple data stream, e.g. SSD

Instructions of x86

SIMD is weakest in case or switch statements, where each execution unit must perform a different operation on its data.

Amdalhs Law

$$\text{Maximum speedup from parallelism} = \frac{1}{((1-P) + \frac{P}{N})}$$

P = proportion of program parallelizable N = number of cores

Signed vs. Unsigned MIPS

sra - shift right arithmetic: sign extension! right shift for signed quantities

srl - shift right logical: zero extension! right shift for unsigned quantities

sll - shift left logical: zero extension! left shift for signed & unsigned quantities

Recursive MIPS Example

Int sum (int n)	sum:
{	addi \$sp, \$sp, -8 # allocate space on stack
return n ? n + sum(n-1): 0;	sw \$ra, 0(\$sp) # store the return address
}	sw \$a0, 4(\$sp)
	# store the argument slti \$t0, \$a0, 1
	# check if n > 0
	beq \$t0, \$0, recurse # n > 0 case
	add \$v0, \$0, \$0 # start return value to 0
	addi \$sp, \$sp, 8 # pop 2 items off stack
	jr \$ra # return to caller
	recurse:
	addi \$a0, \$a0, - 1 # calculate n-1
	jal sum # recursively call sum(n - 1)
	lw \$ra, 0(\$sp) # restore saved return address
	lw \$a0, 4(\$sp) # restore saved argument
	addi \$sp, \$sp, 8 # pop 2 items off stack
	add \$v0, \$a0, \$v0 # calculate n + sum(n - 1)
	jr \$ra # return to caller